

# MONITORING MASSIVE NETWORK METRICS

Vicente De Luca

Sr. Network Engineer at Zendesk

vdeluca@zendesk.com

<https://nethero.org>

# Problem we are trying to solve

Poll all load balancers metrics every 1min including:

- Memory utilization breakdown by linux processes
- Backend pool statistics including breakdown per members

Metric volume:

- 3k pool objects per load balancer
- each pool object contain multiple metrics (28)
- also nested member metrics ( $25 \times N$ )  $N =$  amount of nodes
- total ~562k metrics per polling, per active load balancer

# 1st shot

Not reinvent the wheel by trying SNMP polling:

- Not all metrics available in factory MIBs (Ex. memory breakdown)
- Static creating custom MIBs don't scale (we tried)
- Hit a wall in CPU resources used by net-snmp (daemon crashes)

Feb 12 11:20:13 lb01 emerg logger: Re-starting snmpd

Feb 12 11:21:14 lb01 emerg logger: Re-starting snmpd

Vendor support provides no viable alternative for this scenario

# Discovering alternatives

We realized we could:

- Get all the stats by CLI show cmds without harming CPU
- Cook a parser for extracting names,tags and values
- Use Datadog for our time series DB, dashboard panel and alerts system

How our data looks like ?

```
ltm pool POOL-INOG {
    active-member-cnt 2
    connq-all.age-edm 0
    connq-all.age-ema 0
    connq-all.age-head 0
    connq-all.age-max 0
    connq-all.depth 0
    connq-all.serviced 0
    connq.age-edm 0
    connq.age-ema 0
    connq.age-head 0
    connq.age-max 0
    connq.depth 0
    connq.serviced 0
    cur-sessions 605
    members {
        server1.inog.net:80 {
            addr 10.0.0.2
            serverside.bits-in 1289371
            serverside.bits-out 31293
            serverside.cur-conns 302
            serverside.max-conns 1000
            serverside.pkts-in 31920
            serverside.pkts-out 31289
            serverside.tot-conns 800
            session-status enabled
            status.availability-state available
            status.enabled-state enabled
            status.status-reason Pool member is available
            tot-requests 132913
        }
    ...
}
```

# Why Datadog ?

- Wide utilized by dev/ops allowing easy correlation graphs
- Increase audience on network metrics
- No infra concerns on scaling up the amount of pushed metrics

# Do I need to pay for Datadog ?

No. Similar approach should work with statsd and compatible backends suchs as InfluxDB

## How (dog)statsd works?

- local agent (dogstatsd) listen to UDP messages
- expect metrics in the following format:

metric.name:value|type|@sample\_rate|#tag1:value,tag2

# And now, what ?

We cooked a script (python) to:

- parse the load balancer CLI show output
- extract metric name, tags and values
- craft and send an UDP packet to dogstatsd for each metric

UDP payload example:

```
netops.lb.serverside.cur_conn:143321|g|#pod:1,#netdevice:lb01,#vip:inog,#port:80,#view:public
```

# Challenges while writing the parser (python) script

- balancing curly brackets its not an easy job
- lucky day: our load balancer output looks like JSON
- forced show output to be JSON by regex replace
- result is a python dictionary where for loops can extract name, tags and values

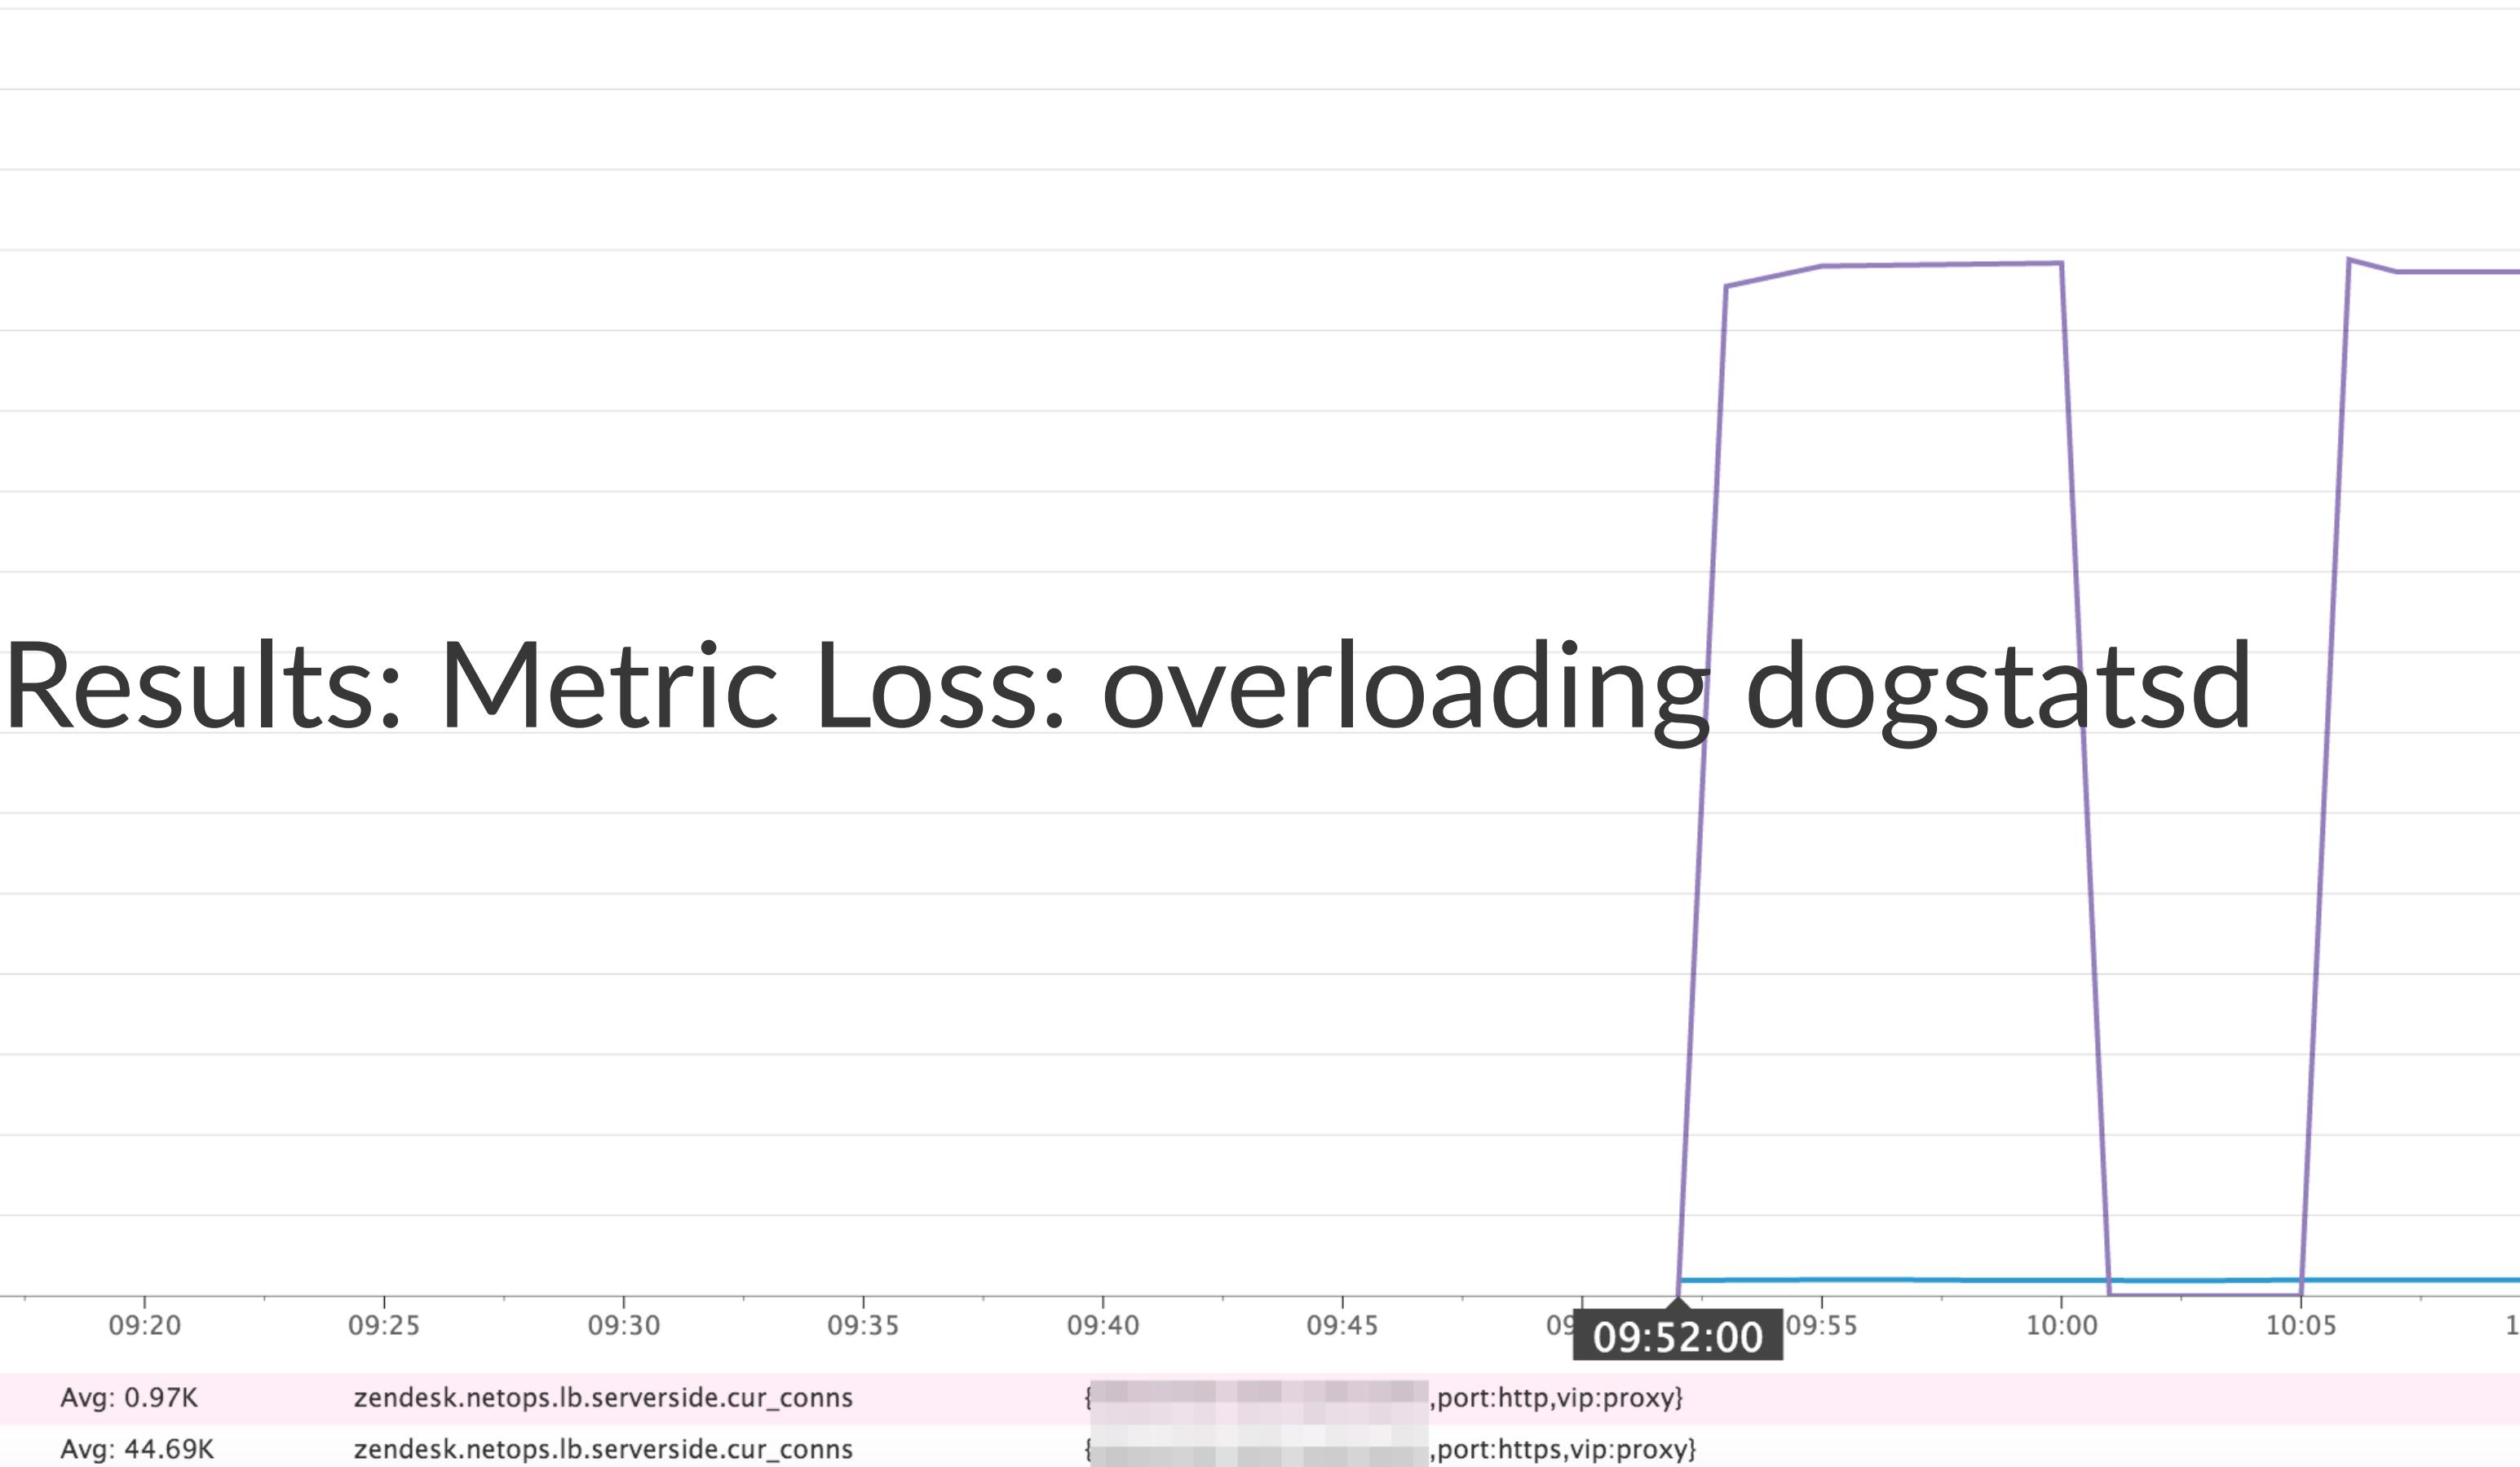
## 2nd shot

At pair of active/standby load balancers:

- bash script execute show cmd every 1 minute, compress the result and send to a linux host via netcat

At linux host:

- nc receives the data, uncompress and call a python script
- python script parse the metrics, extract tags and values
- craft and submit UDP datagrams to local dogstatsd process



## 3rd shot (+Improvements)

At pair of active/standby load balancers:

- bash script now checks if unit is active before submit metrics via netcat (reduced by half workload on dogstatsd)

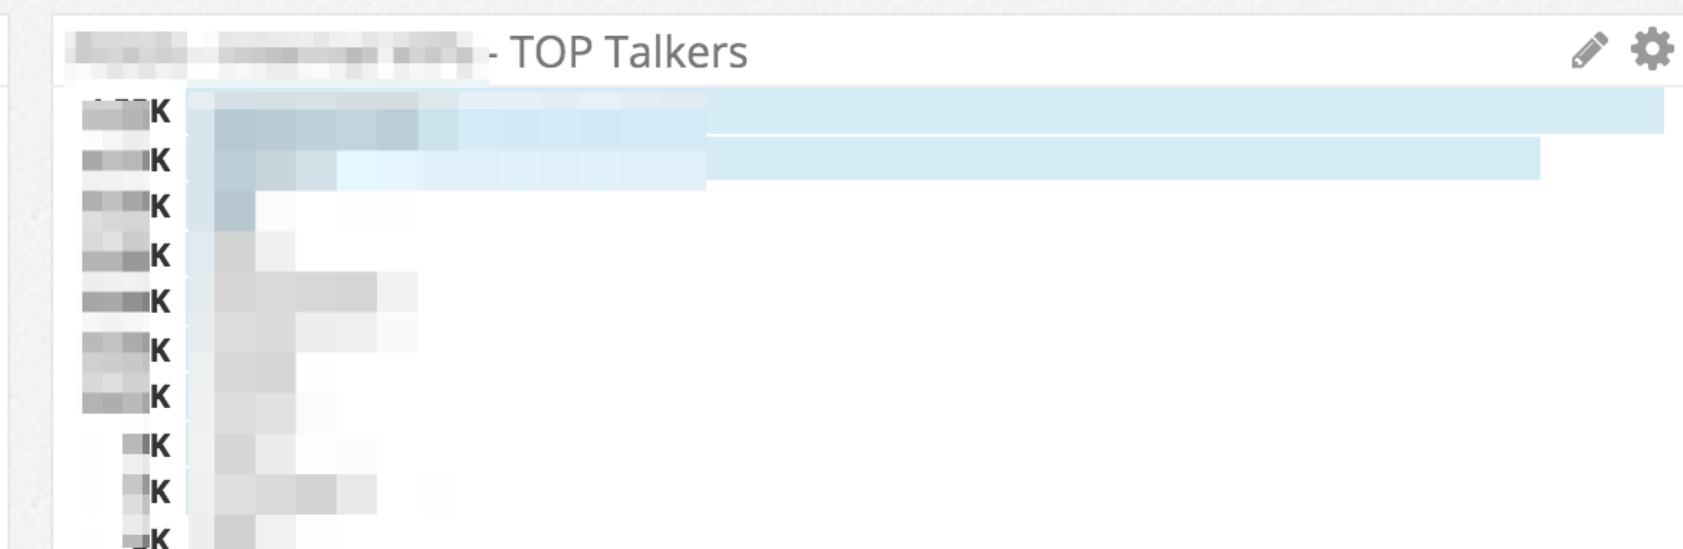
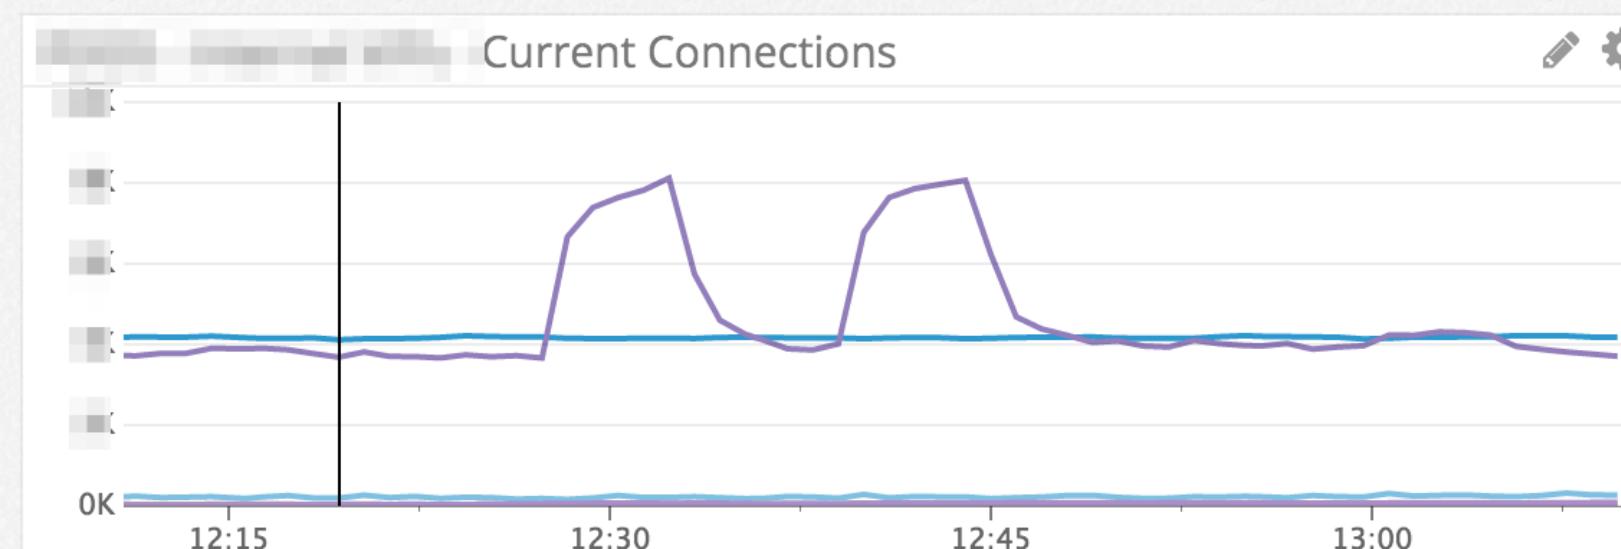
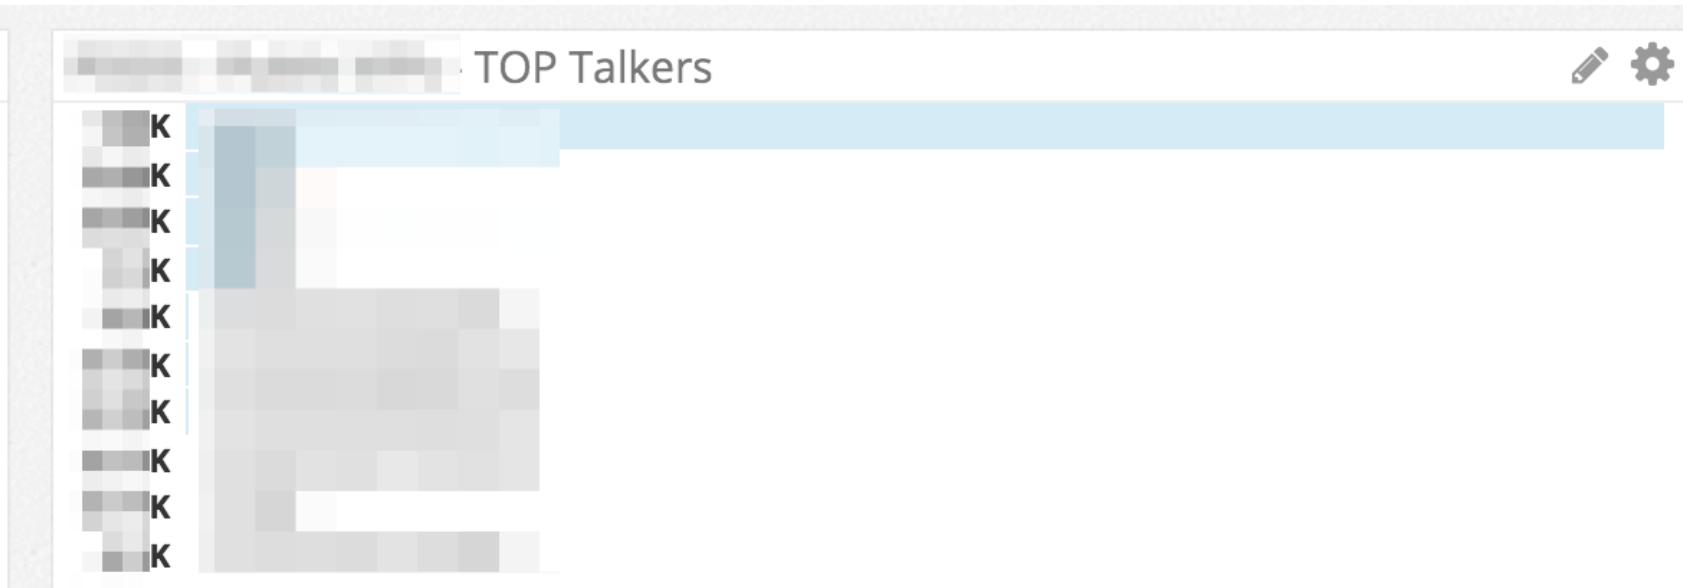
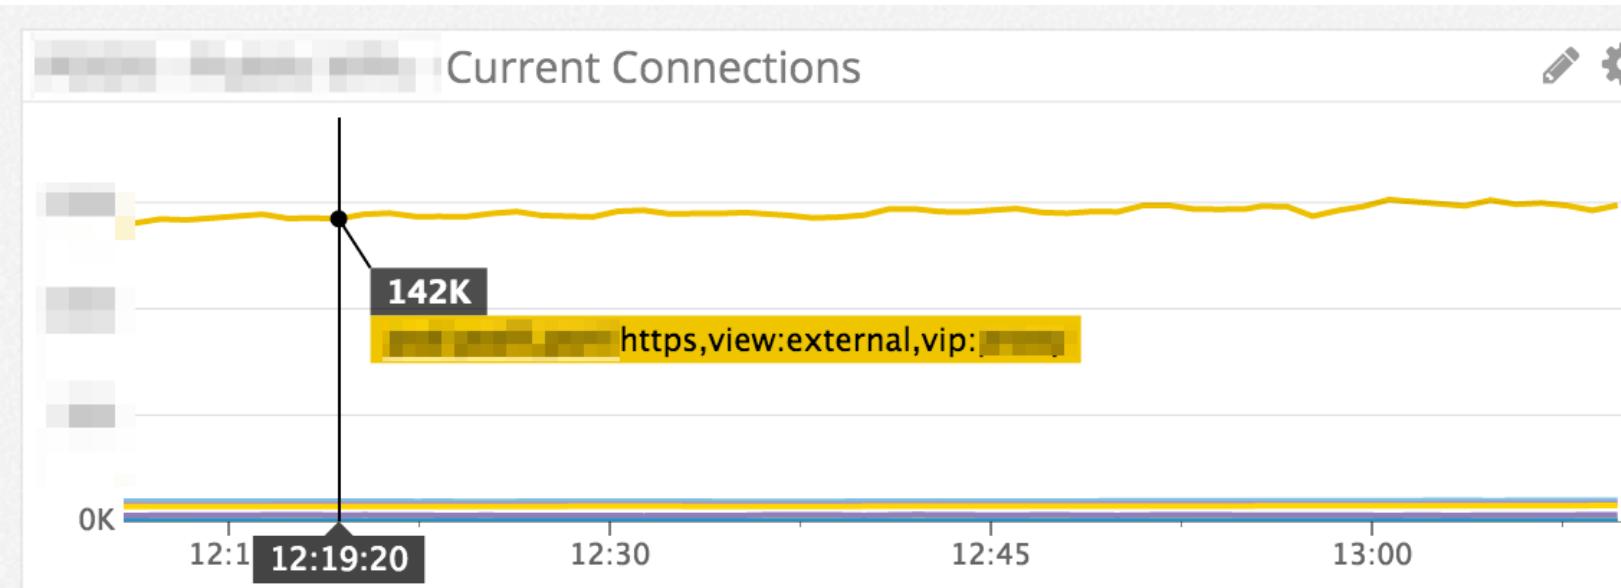
At linux host:

- Filter any non need metrics (all permanent zeroed values)
- Splay: Python send the UDP packets in large blocks, sleeping a few before submitting next block

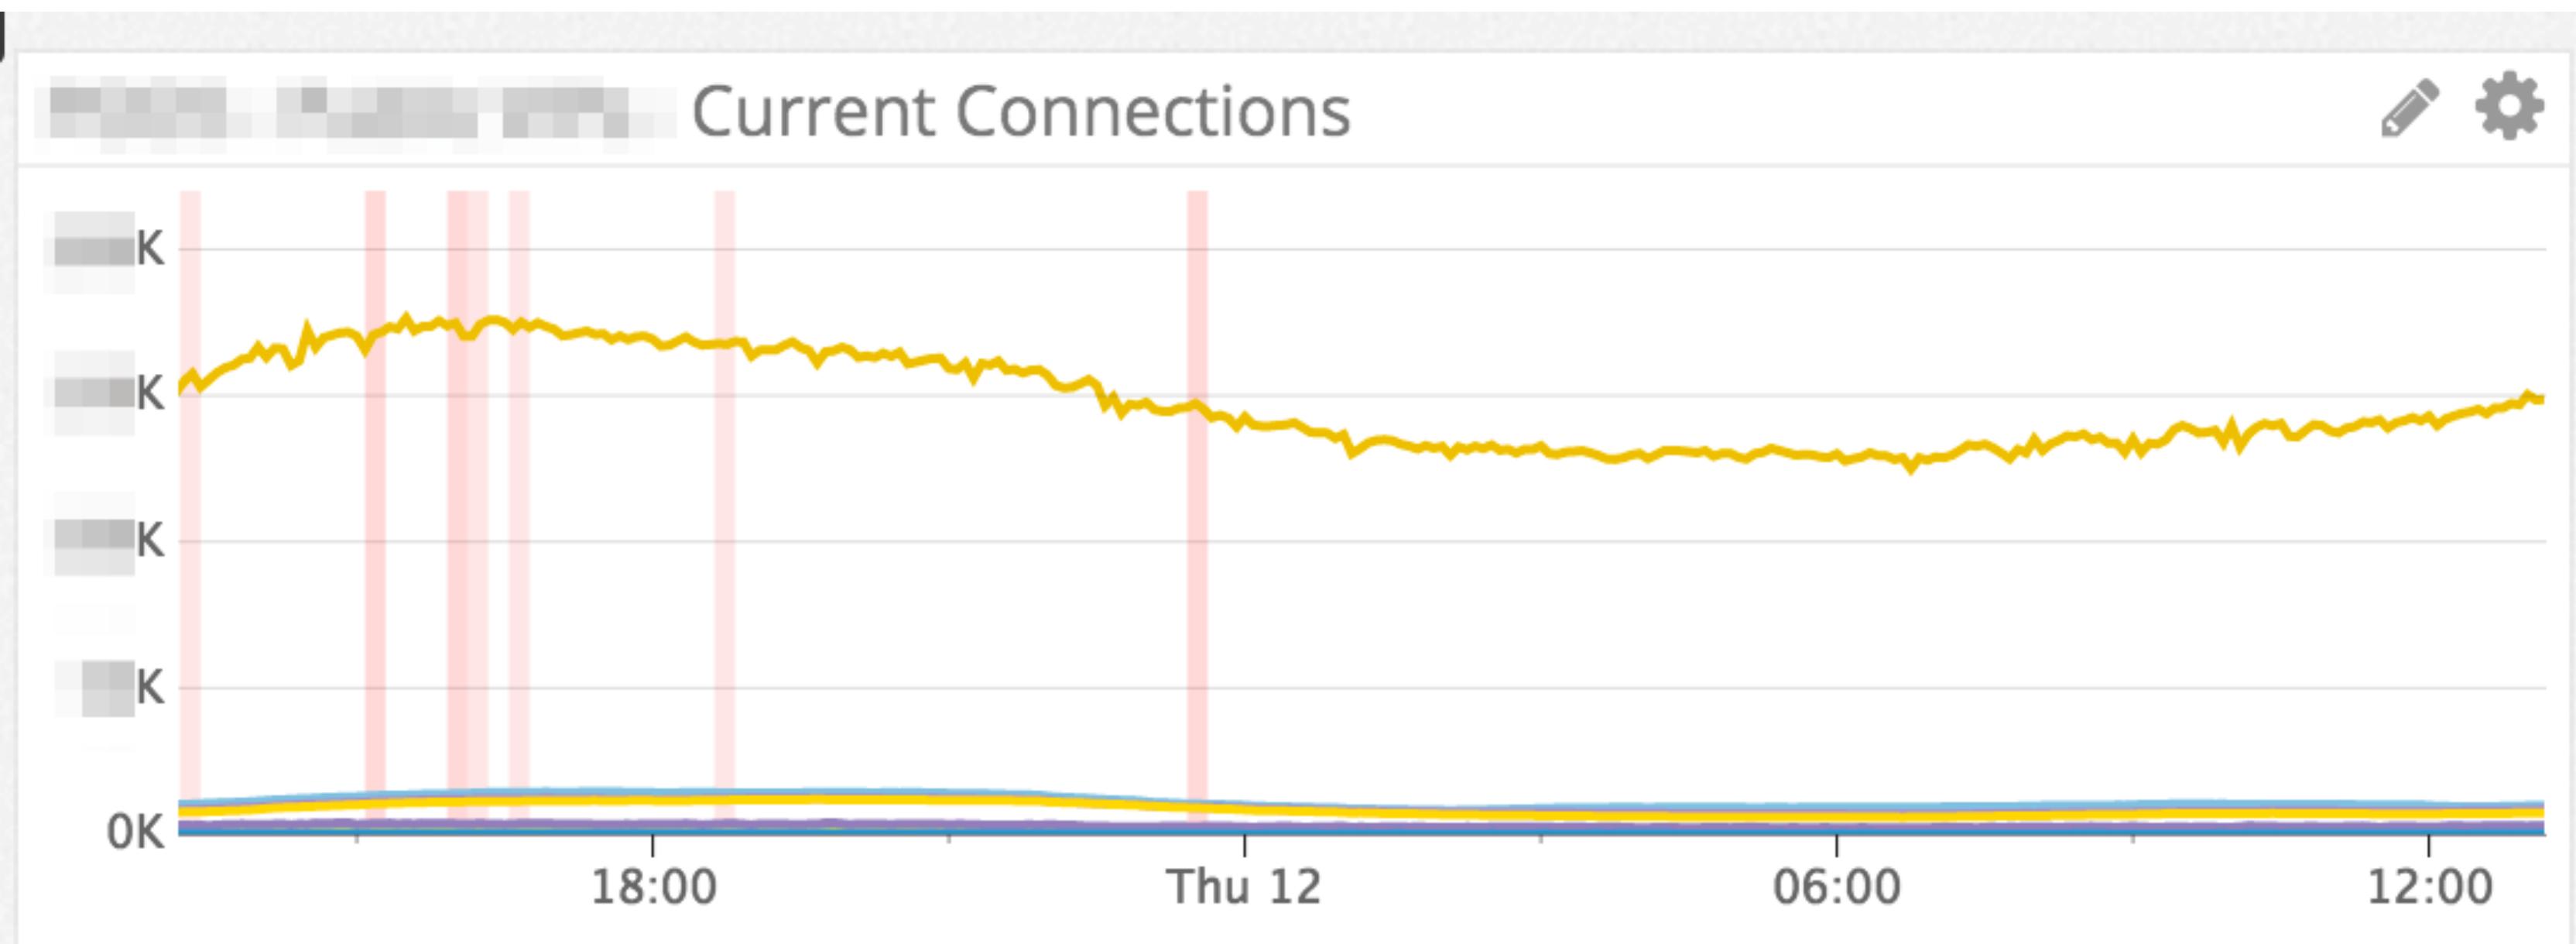
# Results

How our data looks now ?

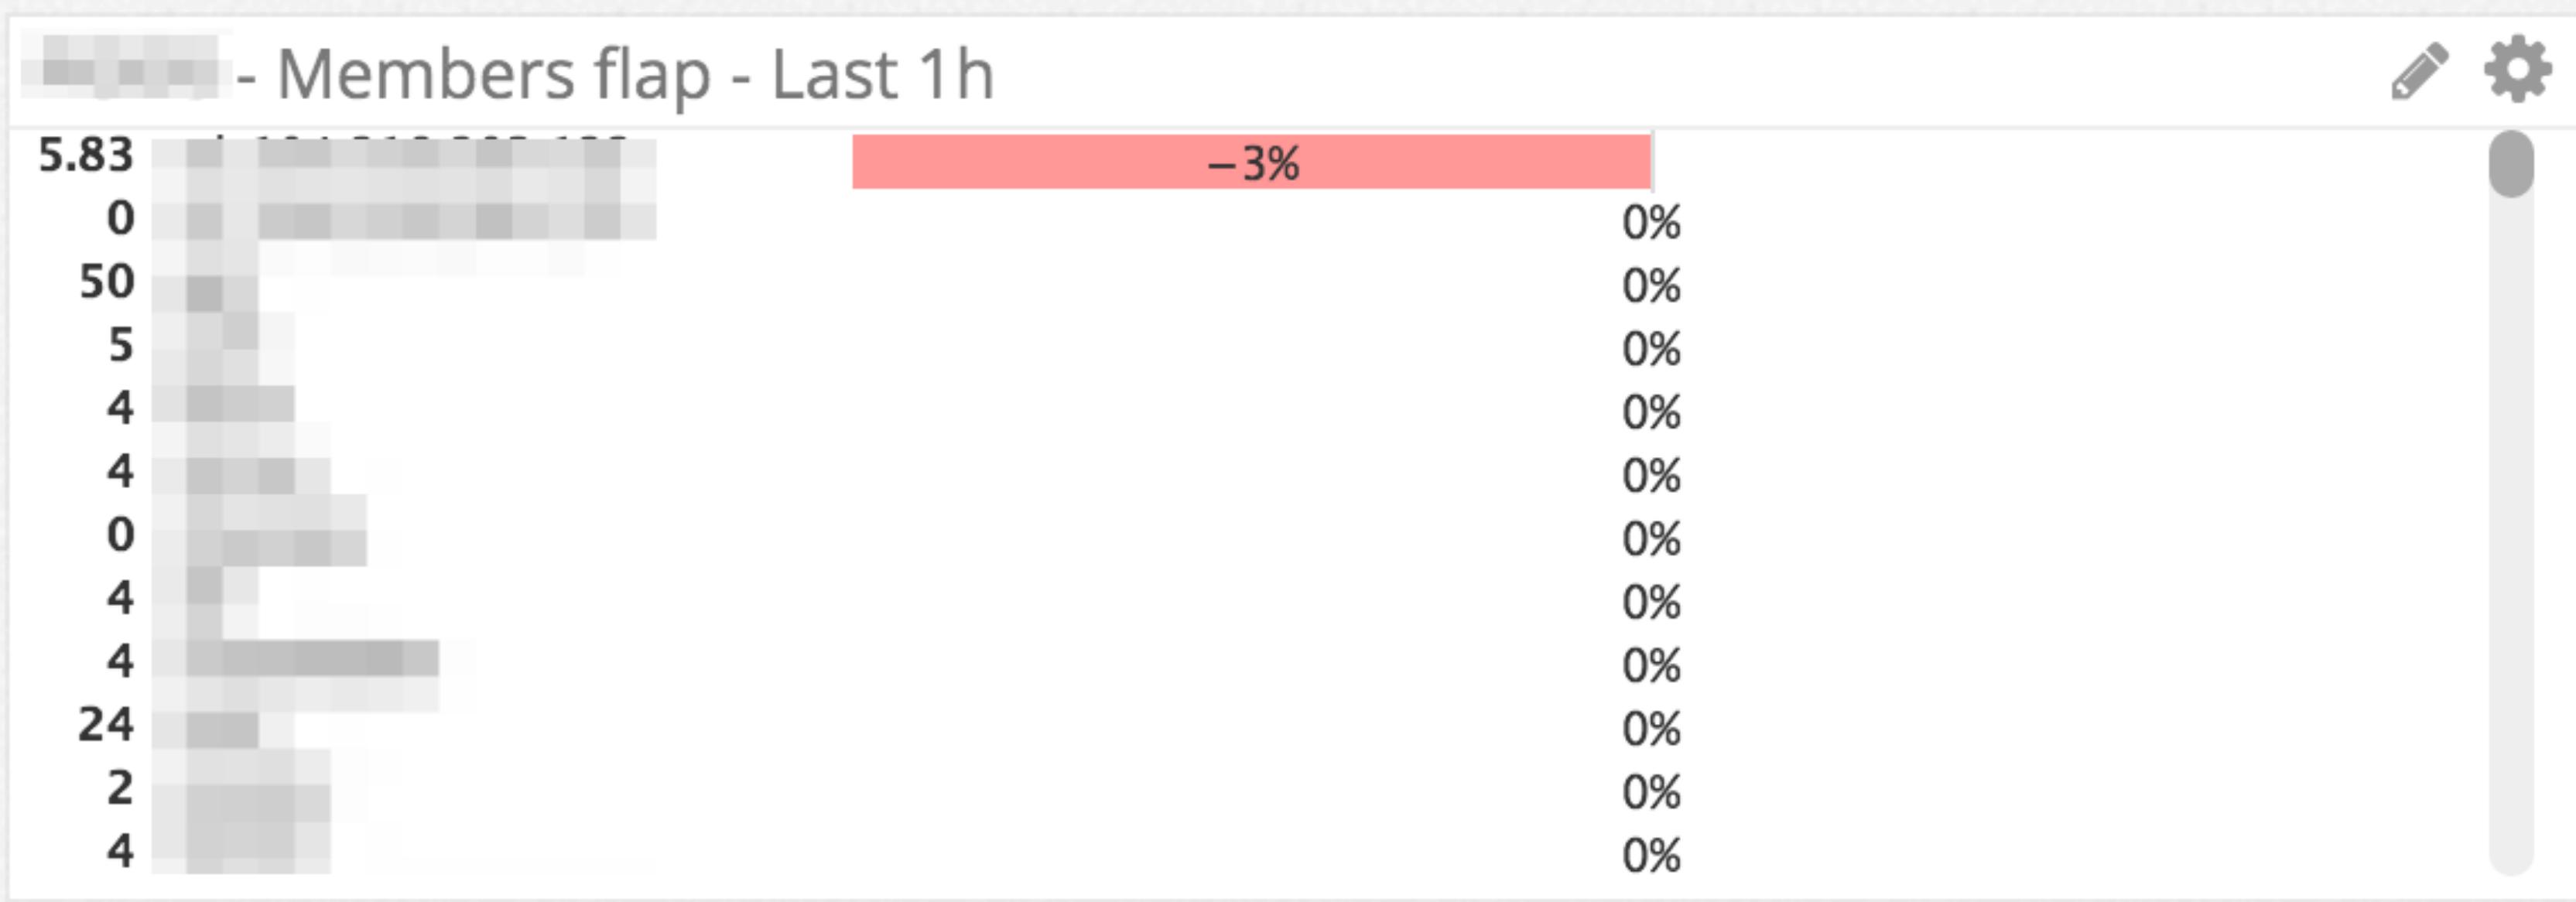
# Current Connections / Top Talkers (every 1 min)



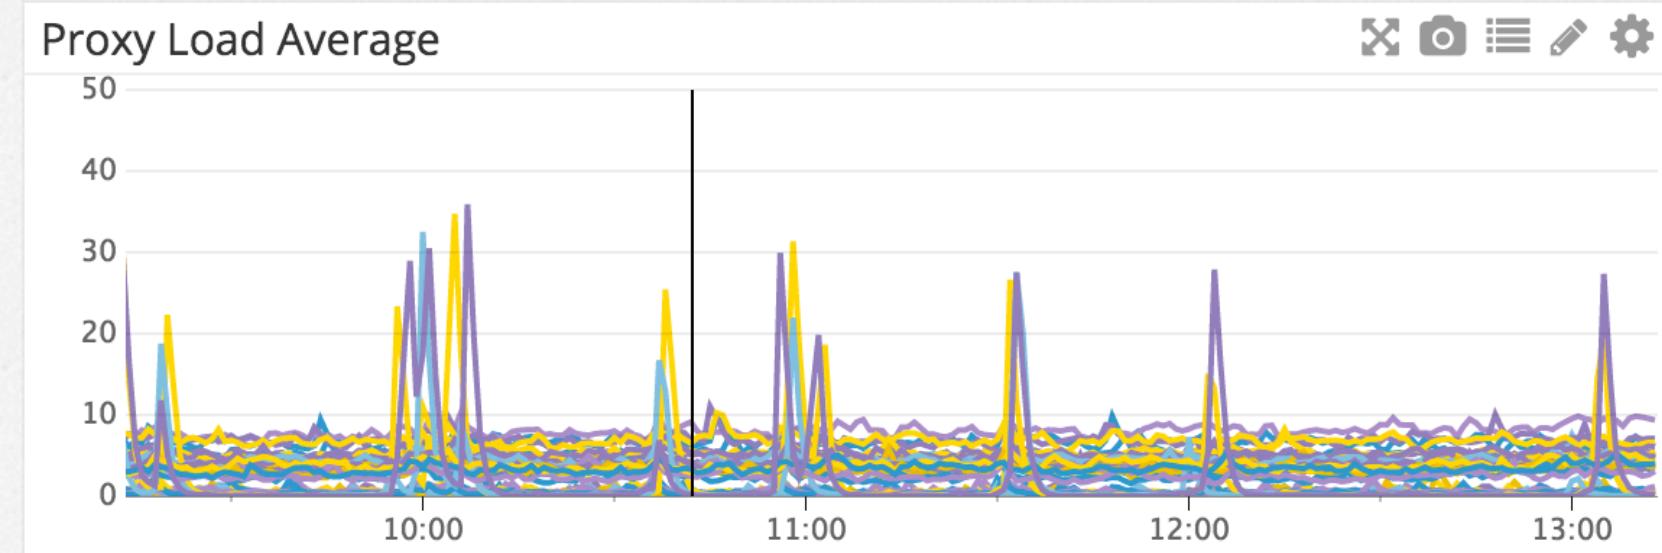
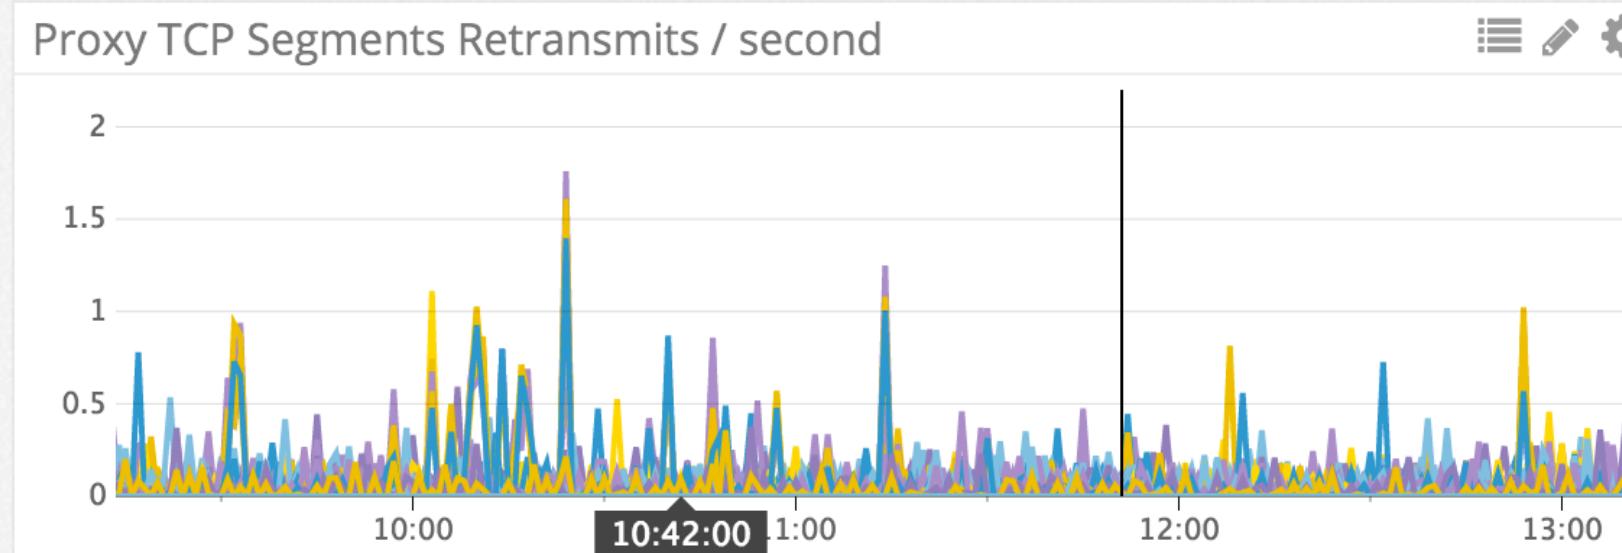
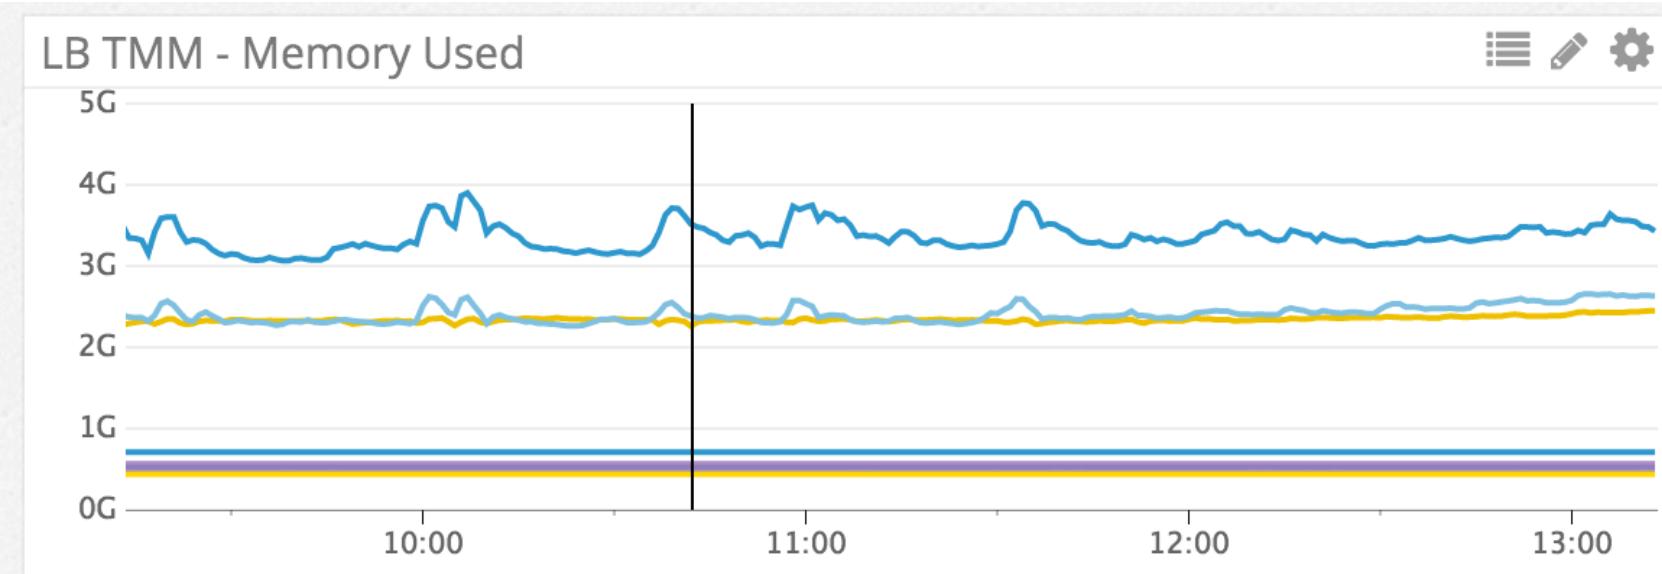
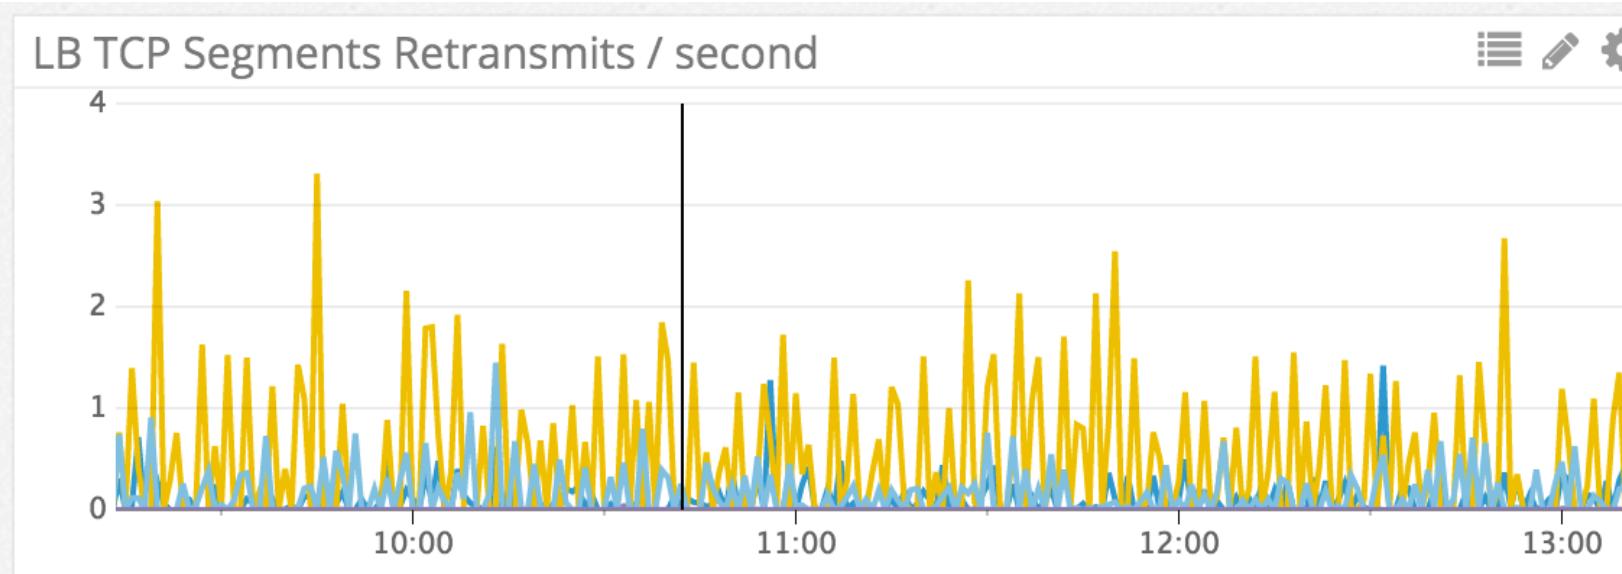
# Events overlay (Network Changes)



# Members flap



# Infrastructure correlation



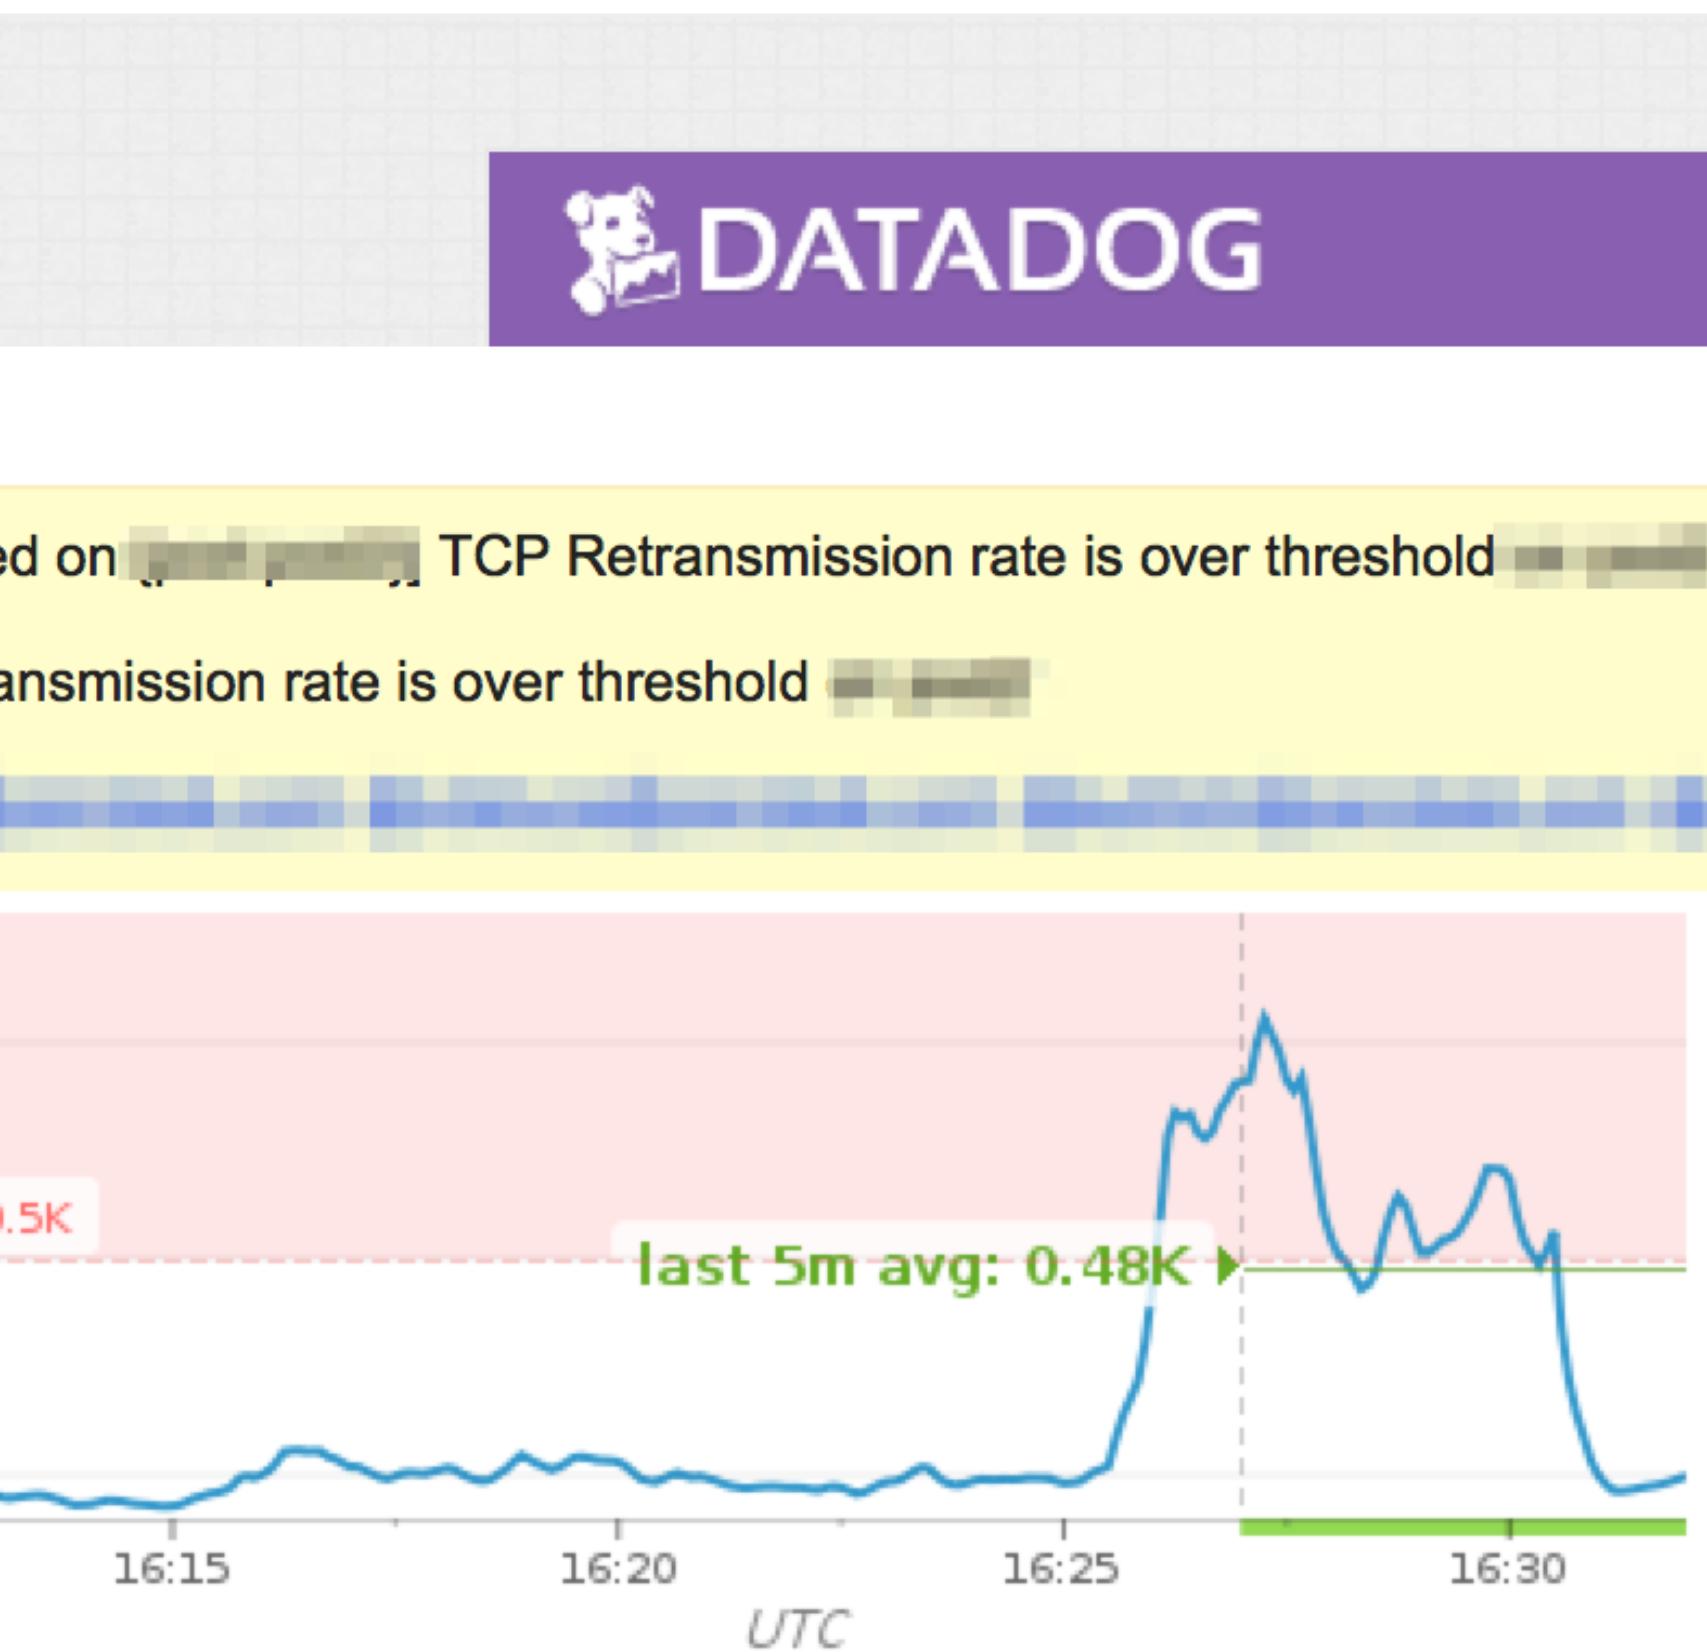
# Results

## Alerting

- Email
- Page on call
- all usual means

Triggered by:

- Configured Thresholds
- Outlier detection algorithms



# Python script execution output

```
2016-05-12 12:56:58.231643 - lb01 - #31383 Metrics processed in 2.79145097733
2016-05-12 12:57:11.526415 - lb01 - #2222 Metrics processed in 0.179049015045
2016-05-12 12:57:26.471943 - lb01 - #2222 Metrics processed in 0.209563970566
2016-05-12 12:57:57.489603 - lb01 - #31383 Metrics processed in 2.79893708229
2016-05-12 12:58:17.208466 - lb01 - #31383 Metrics processed in 2.70802783966
2016-05-12 12:58:30.733715 - lb01 - #2222 Metrics processed in 0.22886300087
2016-05-12 12:58:45.601631 - lb01 - #2222 Metrics processed in 0.184427976608
2016-05-12 12:59:15.377633 - lb01 - #31383 Metrics processed in 2.51854896545
2016-05-12 12:59:28.962007 - lb01 - #2222 Metrics processed in 0.179329872131
```

# Advantages

- Dynamic discovery for new pools
- Easy metric correlation between network, servers or applications
- Anomaly (outliers) detection algorithms
- Derisive CPU consumption compared to SNMP

# QUESTIONS?

Special thanks for contributors here

- Cassiano Aquino ([caquino@zendesk.com](mailto:caquino@zendesk.com))
- Stephen O'Neill ([soneill@zendesk.com](mailto:soneill@zendesk.com))