

Detecting and Mitigating DDoS

A FastNetMon (+more) use case



AS61186 / AS21880



Vicente De Luca

Sr. Network Engineer, ZENDESK



Follow me on Twitter: @zenvdeluca

Follow me on Github: zenvdeluca

Pay me a visit a <http://nethero.org>

vdeluca@zendesk.com

The screenshot shows a desktop application window for managing customer tickets. At the top, there's a green header bar with the text "Changing resolution after the fact" and a yellow "+ add" button. Below the header, the user profile "Jessie Prestige" is shown, along with a "NEW Ticket #40". The main area displays a list of tickets categorized by status: New, Open, Solved, and Closed.

User Profile:

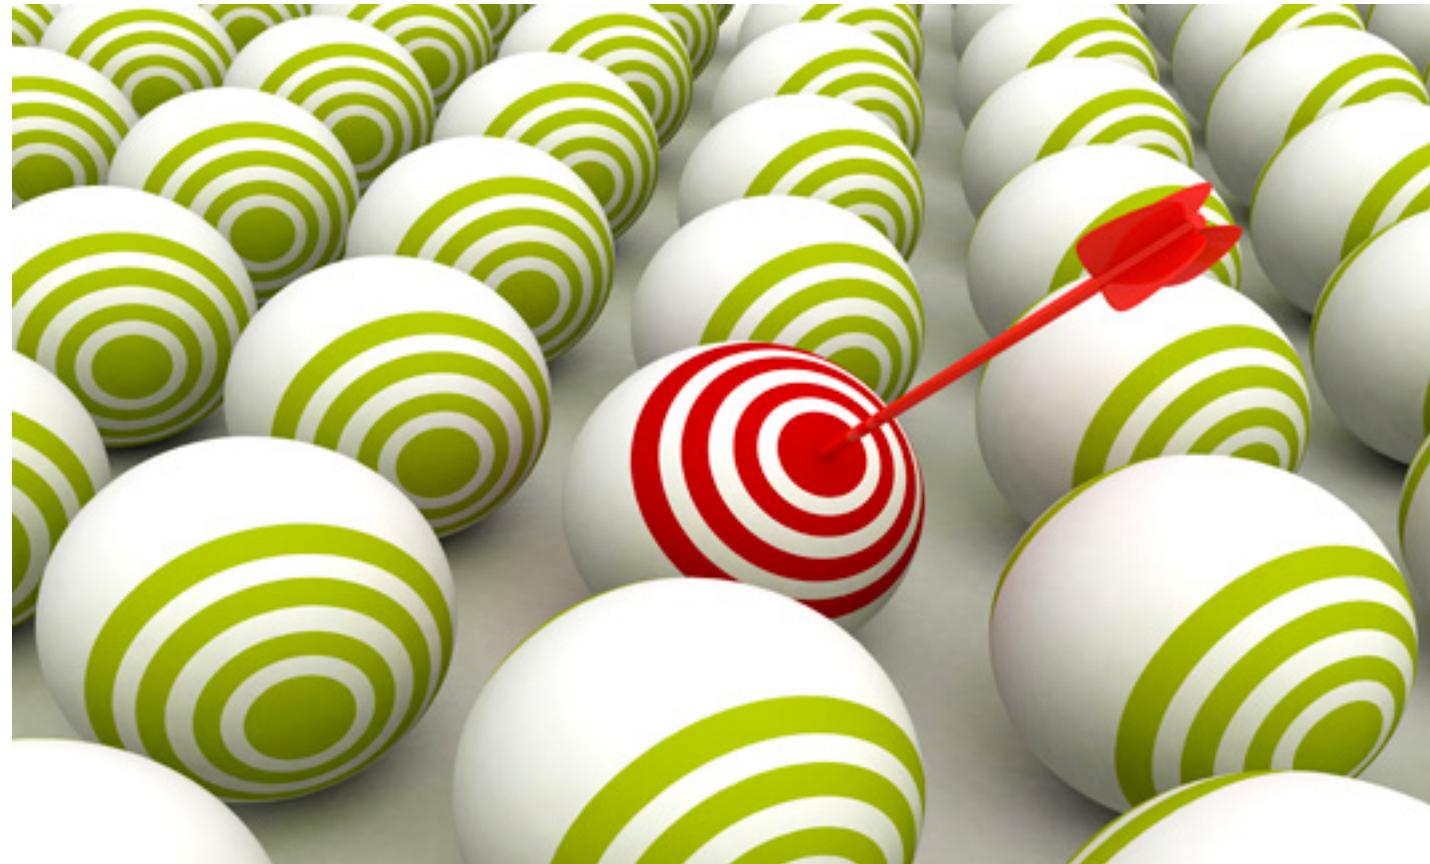
- Role: End-user
- Access: Tickets requested by user
- Email: jessie.prestige@gmail.com
- Twitter: @JessiePrestige
- Phone: 5552752781
- Add contact
- Tags: -
- Org.: Fish Eye
- Language: English

Ticket List:

<input type="checkbox"/>	SUBJECT	REQUESTER	REQUESTED	GROUP
<input type="checkbox"/>	Changing resolution after the fact	Jessie Prestige	Aug 31	Premier
<input type="checkbox"/>	New DSLR chip	Jessie Prestige		
<input type="checkbox"/>	Cannot upload my photos!	Jessie Prestige		
<input type="checkbox"/>	What is your return policy?	Jessie Prestige		
<input type="checkbox"/>	What is your return policy?	Jessie Prestige		



Why are we targets?



support.acme.com

CNAME

acme.zendesk.com



The good, the bad
and the ugly



The good cloud mitigation provider

- six scrubbing centers across the globe
- 3 Tbps of mitigation bandwidth
- BGP: advertise the target /24, receive legit traffic via clean pipe

cons:

- Reaction time: Internet route convergence (BGP)
- mitigation occurs on incoming only
- always on = very pricy



The bad

NOC find out about an ongoing
DDoS attack with site-down alert



The ugly

detecting takes too long
all dependent on humans :(

Improving detection
Automating mitigation

Why not simply use a DDoS mitigation equipment?

- \$\$\$\$ +\$\$\$\$\$\$\$\$!!!
- Need dedicated qualified network engineers
- Demand significant changes in network architecture
- Not simple to automate with modern APIs out in the wild
- Mitigation is useless in case of channel overflow

Open-source recipe

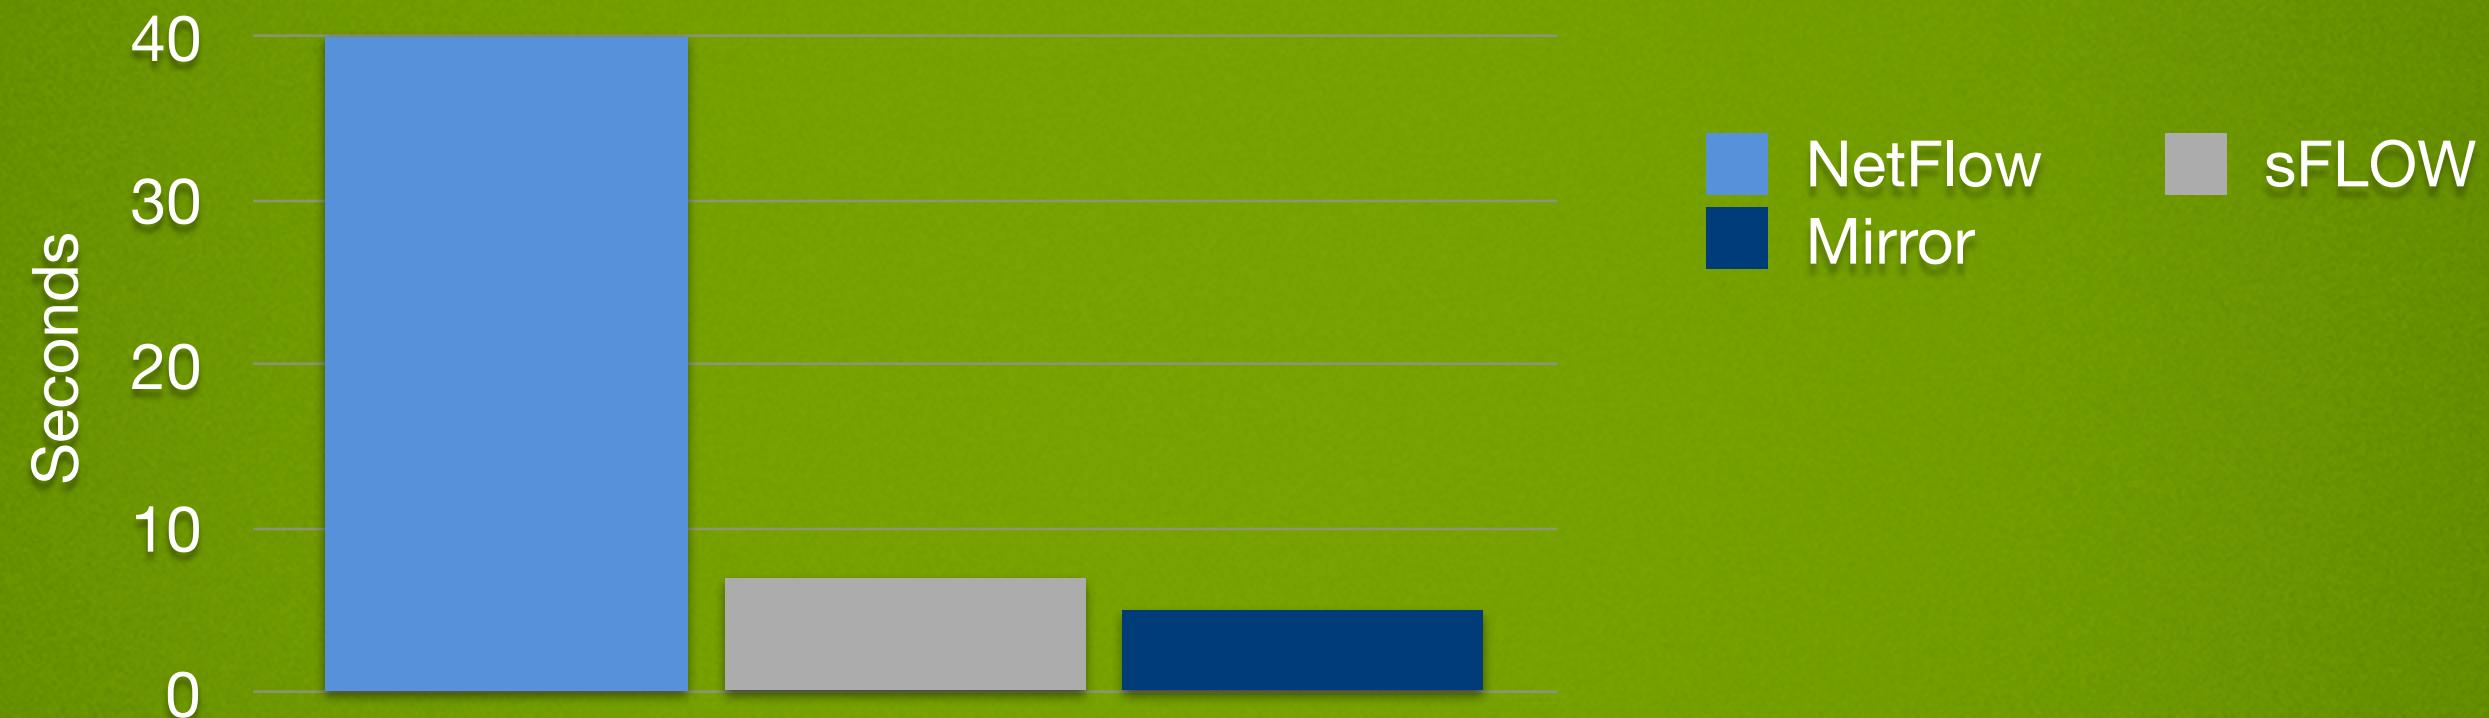
- **FastNetMon:** main core of our solution. DDoS analyzer with sflow/netflow/mirror support
- **InfluxDB:** Scalable data store for metrics, events, and real-time analytics.
- **Grafana:** Gorgeous metric viz, dashboards & editors
- **Redis:** An in-memory database that persists on disk.
- **Morgoth:** Metric anomaly detection for Influx databases
- **BIRD:** a fully functional dynamic IP routing daemon
- **Net Healer:** experimental code to "glue" all moving parts, trigger actions and provide API

How does it work?

FastNetMon - the main core

- supports sFlow (v4/v5), NetFlow (v5/v9/v10), IPFIX and SPAN/mirror
- officially supported in CentOS / Ubuntu / Debian / Vyatta / FreeBSD
- source code available for compiling on your platform
- fast detect IPv4 targets and notify by email, script or populating RedisDB
- BGP support - can advertise routes with communities (for blackhole), or BGP FlowSpec (RFC 5575)
- Tested with Juniper, Cisco, Extreme, Huawei and Linux (ipt_NETFLOW)

FastNetMon: Detection time for capture backends



FastNetMon - the main core

- Works by configured thresholds for mbps, pps or flows:
 - global
 - host (/32)
 - host group (CIDR)
 - per protocol (tcp/udp/icmp)

our in-house tests

Cycle
DDoS Attack started

FNM quiescence:
15s per /32

FastNetMon detect the attack
ban = add /32 attack report
on Redis DB

Net Healer watches Redis DB
if the attack report(s) match a policy, trigger an action

Net Healer Policies example:

(in a time period of 5 min)

if attack reports = 2 then trigger on call
if attack reports ≥ 4 then inject /24 route

if attack report = 1 + anomaly detected (morgoth)
then trigger on call + inject /24 route

time window / policies can be customized

Why Net Healer ?

- FNM relies on pre-configured thresholds
- Hard to predict realistic thresholds since our traffic is influenced by our customers activity (out of our control)
- To avoid false positives we prefer to trigger different actions based on each attack cycle.
- Allows quick integrations like Morgoth x FNM consensus, or API calls such as Pagerduty, etc.

Why InfluxDB ?

- Speaks graphite protocol (compatible with FastNetMon)
- Drop in binary - simple install
- Supports cluster mode - easy to scale

Why Morgoth ?

- Implements non-gaussian algorithm (MGOF) to detect anomaly on data stream metrics
- Takes InfluxDB (bps/pps) fingerprints every chunk of 10s
- Compares the actual fingerprint with the past learned traffic
- Anomaly detected? create an alert entry and a graph vertical markdown

Why BIRD ?

- synced with linux kernel routing table
- if a route is added on a separate linux routing table, BIRD will learn this route and advertise it to my edge routers.
RPL announces the prefix to DDoS scrubbing center
- friendly to Network Engineers (birdc)

How does it look ?



Dashboards

Zoom Out

⌚ an hour ago to a few seconds ago refreshed every 5s ▾



Data Sources

⚡ Attack Warning ⚡ Attack Critical ⚡ Anomaly bps ⚡ Anomaly pps 

Vicente De Luca



Zendesk



Grafana admin



Sign out

IAD1 - Traffic Bandwidth

in / out (bps ratio)



	max	avg	current
--	-----	-----	---------

total.mean {direction: incoming}	5.06 Gbps	266 Mbps	108 Mbps
total.mean {direction: outgoing}	1.69 Gbps	417 Mbps	954 Mbps
5-incoming	32 Mbps	7 Mbps	8 bps
-outgoing	124 Mbps	31 Mbps	58 Mbps
incoming	53 Mbps	8 Mbps	8 bps
outgoing	123 Mbps	38 Mbps	69 Mbps
ssl-incoming	101 Mbps	808 Kbps	8 bps
ssl-outgoing	9 Mbps	2 Mbps	3 Mbps

0.958

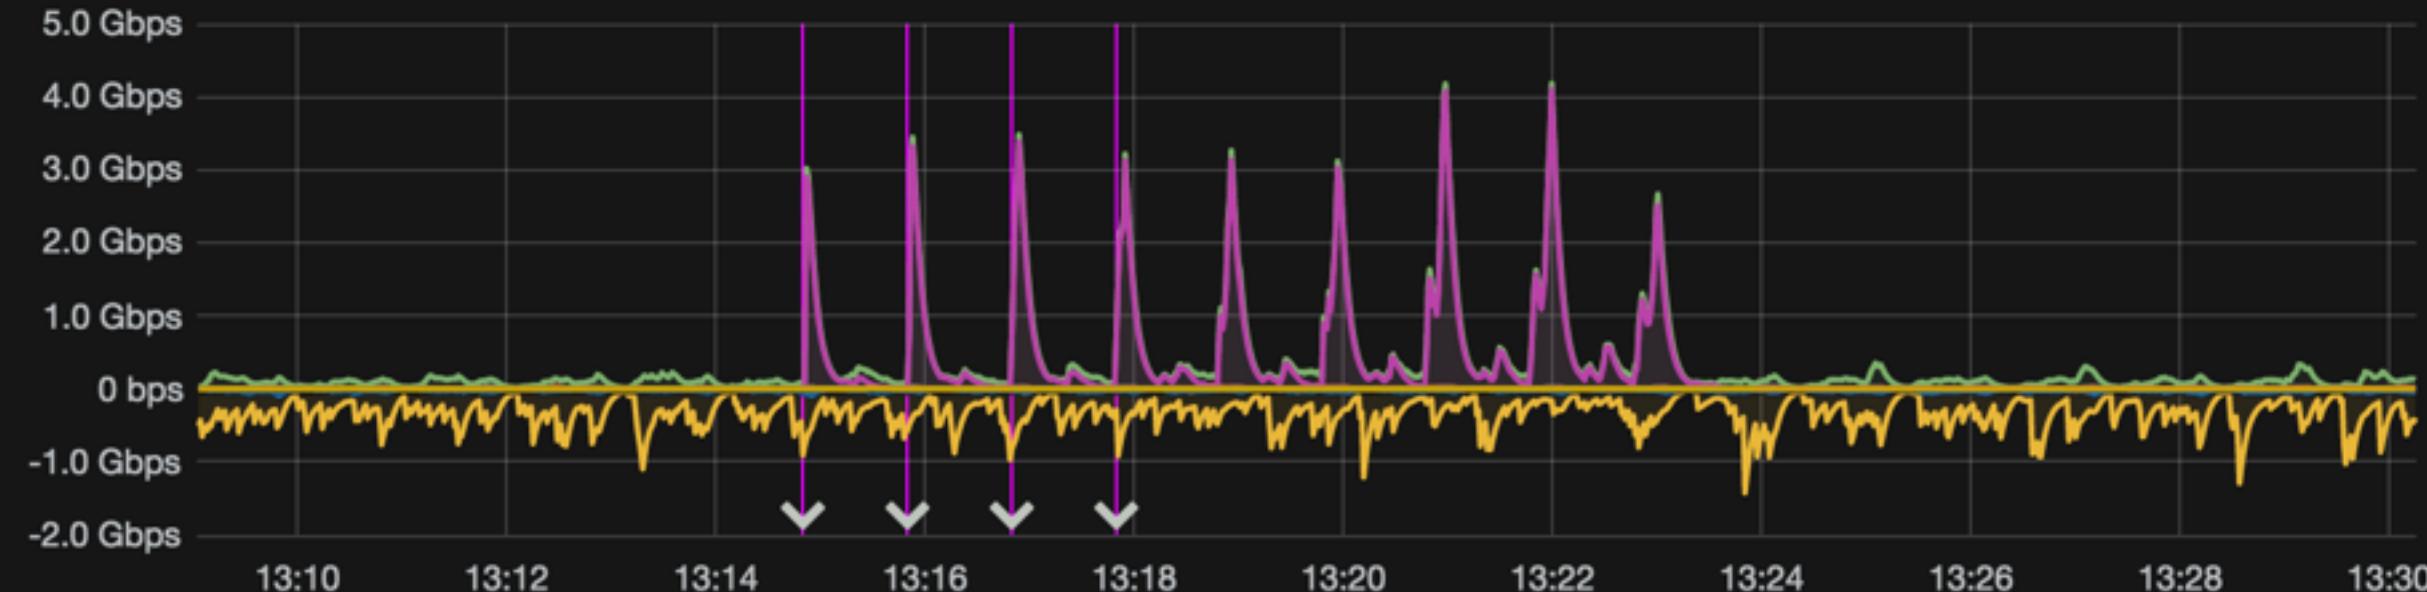
Attack Warning ✓

Attack Critical ✓

Anomaly bps ✓

Anomaly pps ✓

IAD1 - Traffic Bandwidth



in / out (bps ratio)

1.171

	max	avg	current
total.mean {direction: incoming}	4.177 Gbps	328 Mbps	121 Mbps
total.mean {direction: outgoing}	1.425 Gbps	342 Mbps	397 Mbps
total.incoming	22 Mbps	7 Mbps	14 Mbps
total.outgoing	59 Mbps	22 Mbps	36 Mbps
ipsec-incoming	53 Mbps	9 Mbps	14 Mbps
ipsec-outgoing	89 Mbps	31 Mbps	49 Mbps
tcp-incoming	4.089 Gbps	323 Mbps	8 bps
tcp-outgoing	7 Mbps	2 Mbps	5 Mbps
http-incoming	4 Mbps	48 Kbps	8 bps
http-outgoing	13 Mbps	144 Kbps	83 Kbps

1 - Packets per second



total-incoming

total-outgoing

tcp-incoming

tcp-outgoing

ssl-incoming

ssl-outgoing

http-incoming

http-outgoing

130 Kpps

194 Kpps

9 Kpps

14 Kpps

9 Kpps

11 Kpps

337 pps

252 pps

15 pps

17 pps

79 Kpps

129 Kpps

5 Kpps

6 Kpps

5 Kpps

7 Kpps

107 pps

112 pps

4 pps

5 pps

113 Kpps

166 Kpps

8 Kpps

8 Kpps

9 Kpps

11 Kpps

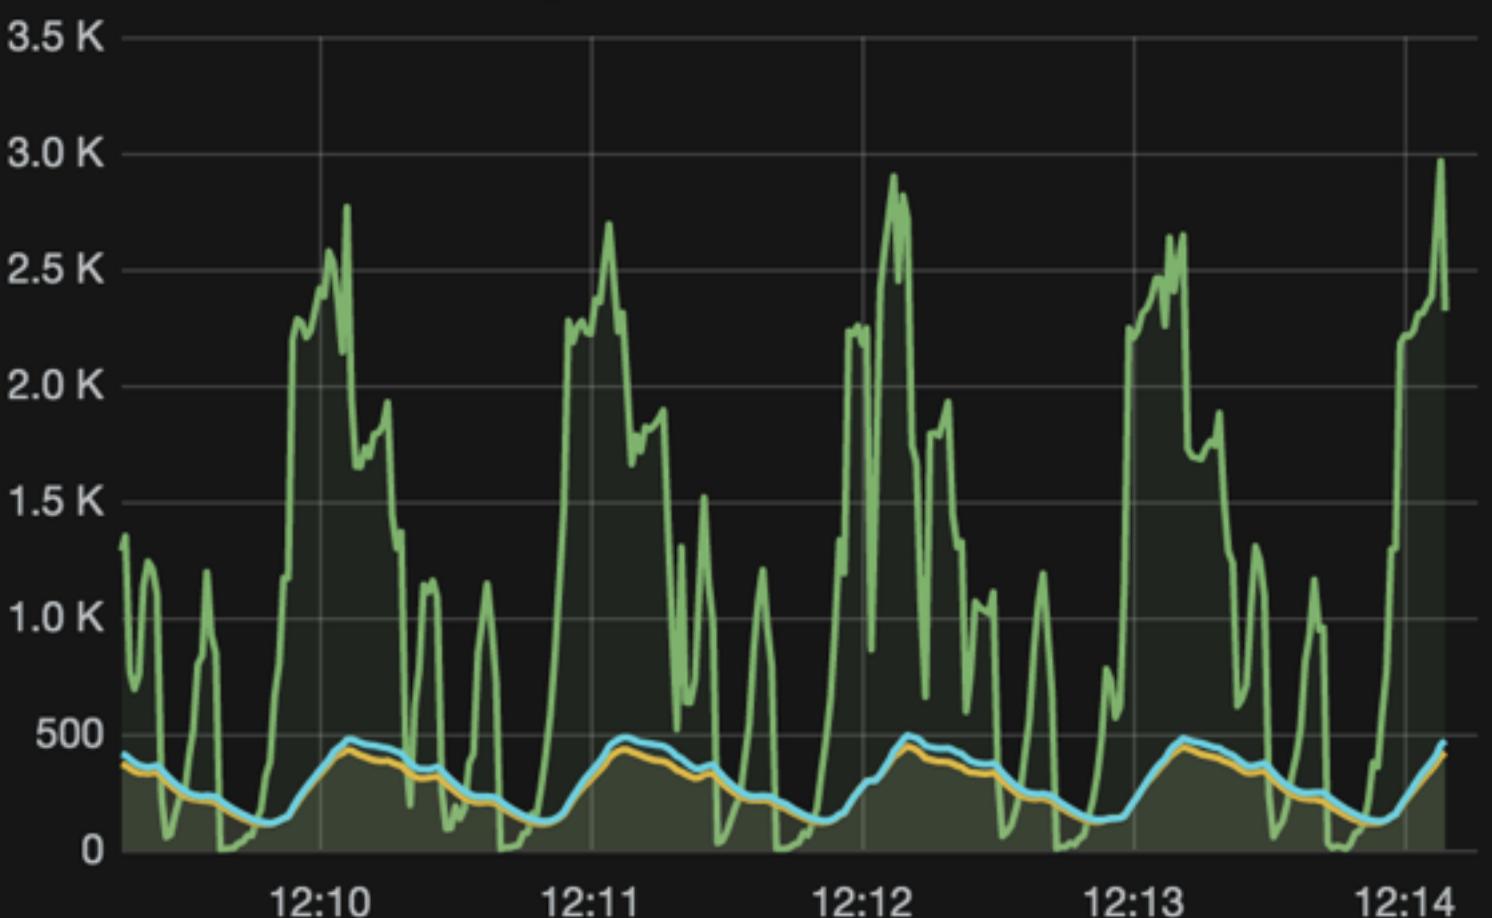
48 pps

70 pps

1 pps

1 pps

1 - Flow amount



total-incoming

tcp-incoming

ssl-incoming

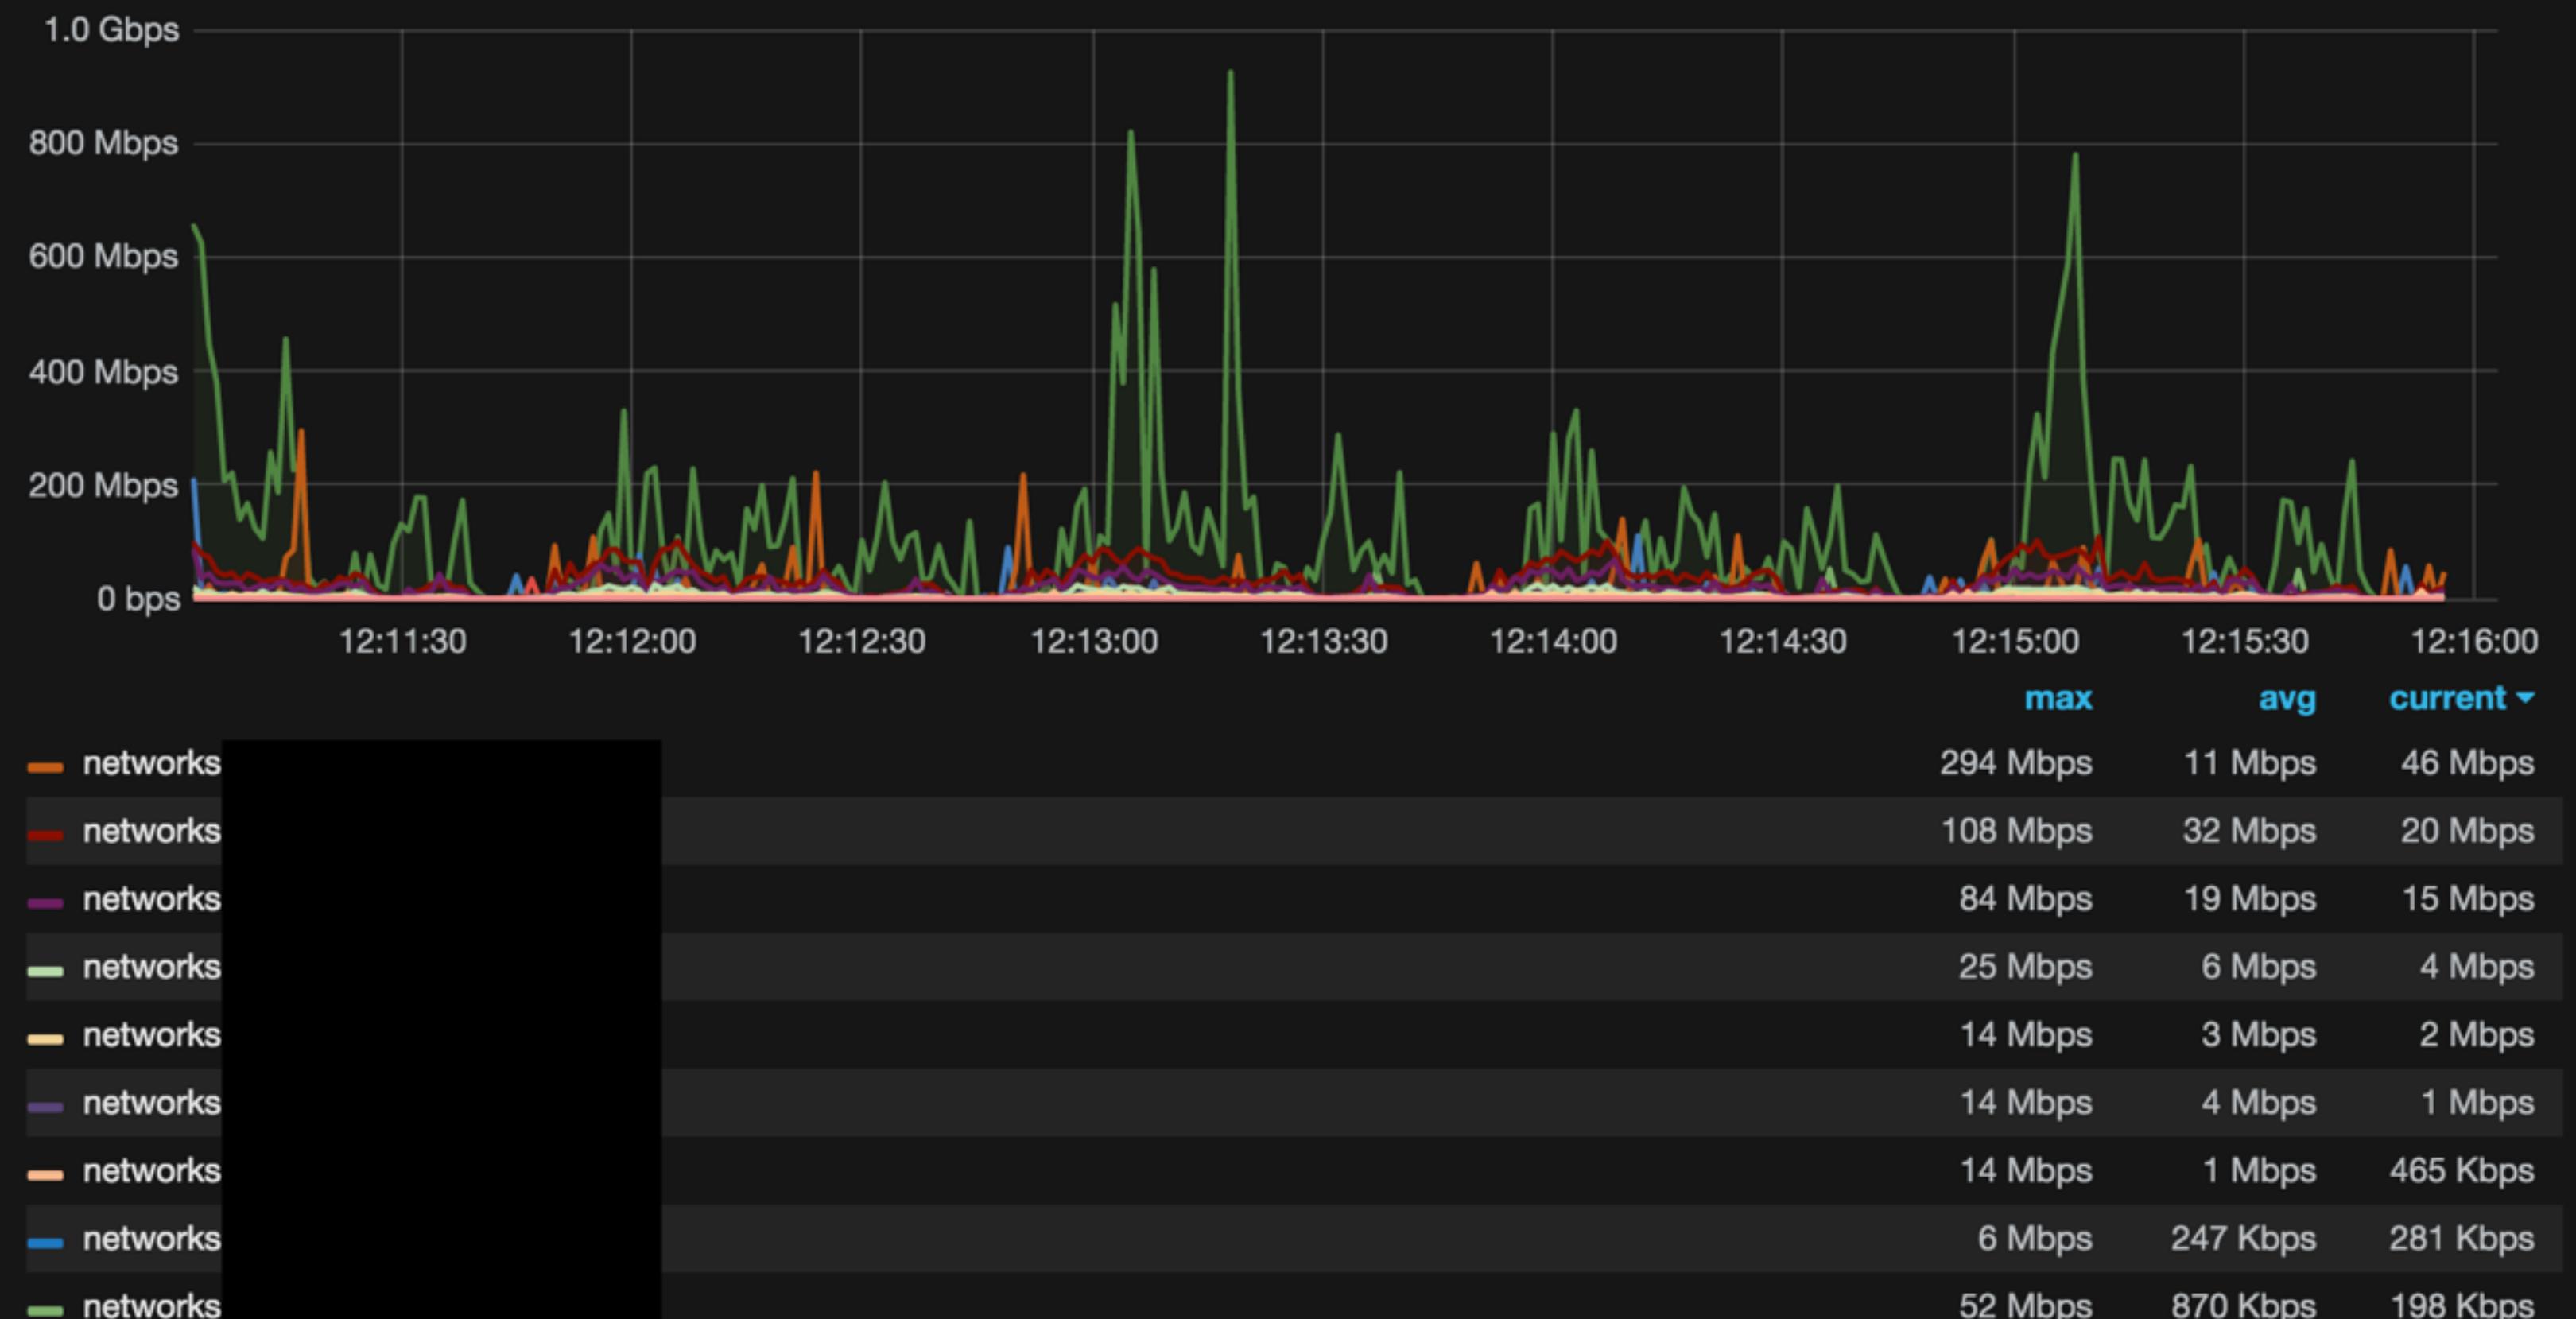
http-incoming

2.963 K

445

492

/24 breakdown - Incoming bps



```
~$> jq . <<< $(curl -sk https://nethealer1.i... /healer/v1/ddos/status)
{
  "status": "clear",
  "timestamp": "20150816-195527"
}
~$> jq . <<< $(curl -sk https://nethealer1.i... /healer/v1/ddos/status)
{
  "status": "warning",
  "target": {
    "192.168.1.1": 3,
    "192.168.1.2": 3
  },
  "timestamp": "20150816-195703"
}
```

```
~$> jq . <<< $(curl -sk https://nethealer1.i... /healer/v1/ddos/status)
{
  "status": "critical",
  "target": {
    "192.168.1.1": 5,
    "192.168.1.2": 5
  },
  "timestamp": "20150816-195926"
}
```

```
~$> jq . <<< $(curl -sk https://nethealer1.10.10.10:443/healer/v1/ddos/reports)
{
  "reports": [
    {
      "ip": "192.168.1.1",
      "attack_details": {
        "attack_type": "unknown",
        "initial_attack_power": 5076,
        "peak_attack_power": 5076,
        "attack_direction": "outgoing",
        "attack_protocol": "tcp",
        "total_incoming_traffic": 1397974,
        "total_outgoing_traffic": 3427164,
        "total_incoming_pps": 3885,
        "total_outgoing_pps": 5076,
        "total_incoming_flows": 210,
        "total_outgoing_flows": 161,
        "average_incoming_traffic": 1397974,
        "average_outgoing_traffic": 3427164,
        "average_incoming_pps": 3885,
        "average_outgoing_pps": 5076,
        "average_incoming_flows": 210,
        "average_outgoing_flows": 161,
        "incoming_ip_fragmented_traffic": 0,
        "outgoing_ip_fragmented_traffic": 0,
        "incoming_ip_fragmented_pps": 0,
        "outgoing_ip_fragmented_pps": 0,
        "incoming_tcp_traffic": 2789304,
        "outgoing_tcp_traffic": 9955449,
        "incoming_tcp_pps": 7817,
        "outgoing_tcp_pps": 13842,
        "incoming_syn_tcp_traffic": 634368,
        "outgoing_syn_tcp_traffic": 1976571,
        "incoming_syn_tcp_pps": 2260,
        "outgoing_syn_tcp_pps": 3225,
        "incoming_udp_traffic": 0,
        "outgoing_udp_traffic": 0
      }
    }
  ]
}
```



From: DDoS Detection
To: Network Operations

Healer Dashboard: <https://netmonitor>.

Attack info:

--
192. [REDACTED].1:

fqdn:

site: IAD1

alerts: 2

protocol:

- udp

incoming:

total:

mbps: 263.59

pps: 628056

flows: 523

tcp:

mbps: 0.9

pps: 2790

syn:

mbps: 0.14

pps: 808

udp:

mbps: 646.55

pps: 9622050

icmp:

mbps: 1.98

pps: 36764

attack_type: udp_flood



- 'UDP flows: 76'

- .1:29126 < 193.15.55.5:123 256392 bytes 2374 packets
- .1:29126 < 83.234.89.5:123 3975192 bytes 9998 packets
- .1:29126 < 37.59.44.6:123 5301036 bytes 11327 packets
- .1:29126 < 46.105.191.10:123 3513420 bytes 7591 packets
- .1:29126 < 217.198.8.13:123 3734928 bytes 9786 packets
- .1:29126 < 109.201.64.18:123 116312976 bytes 248532 packets
- .1:29126 < 5.135.125.18:123 147996 bytes 329 packets
- .1:29126 < 94.190.64.21:123 445680 bytes 2476 packets
- .1:29126 < 185.70.135.26:123 116703756 bytes 249367 packets
- .1:29126 < 95.188.64.32:123 4507272 bytes 10014 packets
- .1:29126 < 198.27.65.33:123 149256 bytes 336 packets
- .1:29126 < 46.105.49.36:123 1235196 bytes 2685 packets
- .1:29126 < 198.91.51.37:123 5482692 bytes 12473 packets
- .1:29126 < 37.187.90.41:123 308880 bytes 660 packets
- .1:29126 < 77.37.160.44:123 452520 bytes 2514 packets
- .1:29126 < 94.23.217.46:123 2492100 bytes 5325 packets
- .1:29126 < 176.31.34.50:123 64620 bytes 199 packets
- .1:29126 < 86.59.35.50:123 7455996 bytes 17435 packets
- .1:29126 < 46.105.74.56:123 235512 bytes 512 packets
- .1:29126 < 85.118.60.57:123 76747788 bytes 163991 packets
- .1:29126 < 5.39.65.67:123 202032 bytes 508 packets
- .1:29126 < 92.63.149.74:123 267840 bytes 2480 packets
- .1:29126 < 149.202.30.78:123 410472 bytes 1084 packets
- .1:29126 < 194.177.31.81:123 619668 bytes 2459 packets
- .1:29126 < 178.47.130.82:123 4165956 bytes 10055 packets
- .1:29126 < 89.203.137.86:123 7477560 bytes 16216 packets
- .1:29126 < 185.50.214.100:123 265896 bytes 2462 packets
- .1:29126 < 208.92.219.100:123 999504 bytes 2524 packets
- .1:29126 < 208.92.219.103:123 3187368 bytes 7582 packets
- .1:29126 < 212.232.61.120:123 2802780 bytes 7523 packets
- .1:29126 < 74.118.17.124:123 112431852 bytes 240239 packets
- .1:29126 < 213.155.226.132:123 1996740 bytes 5037 packets
- .1:29126 < 142.54.178.138:123 1797876 bytes 4983 packets
- .1:29126 < 176.31.104.142:123 190296 bytes 502 packets
- .1:29126 < 5.39.70.150:123 4489992 bytes 9594 packets

Work in progress

all the ingredients on this recipe are open source

About Net Healer: experimental Ruby code ideas, issues and pull requests are **more** than **welcome**

Join FastNetMon mail list

- <https://groups.google.com/forum/#!forum/fastnetmon>

Contribute at Net Healer github

- https://github.com/zenvdeluca/net_healer



Thank you !

vdeluca@zendesk.com